



Проблемы PQ

Андрей Бородин, Postgres Contributor

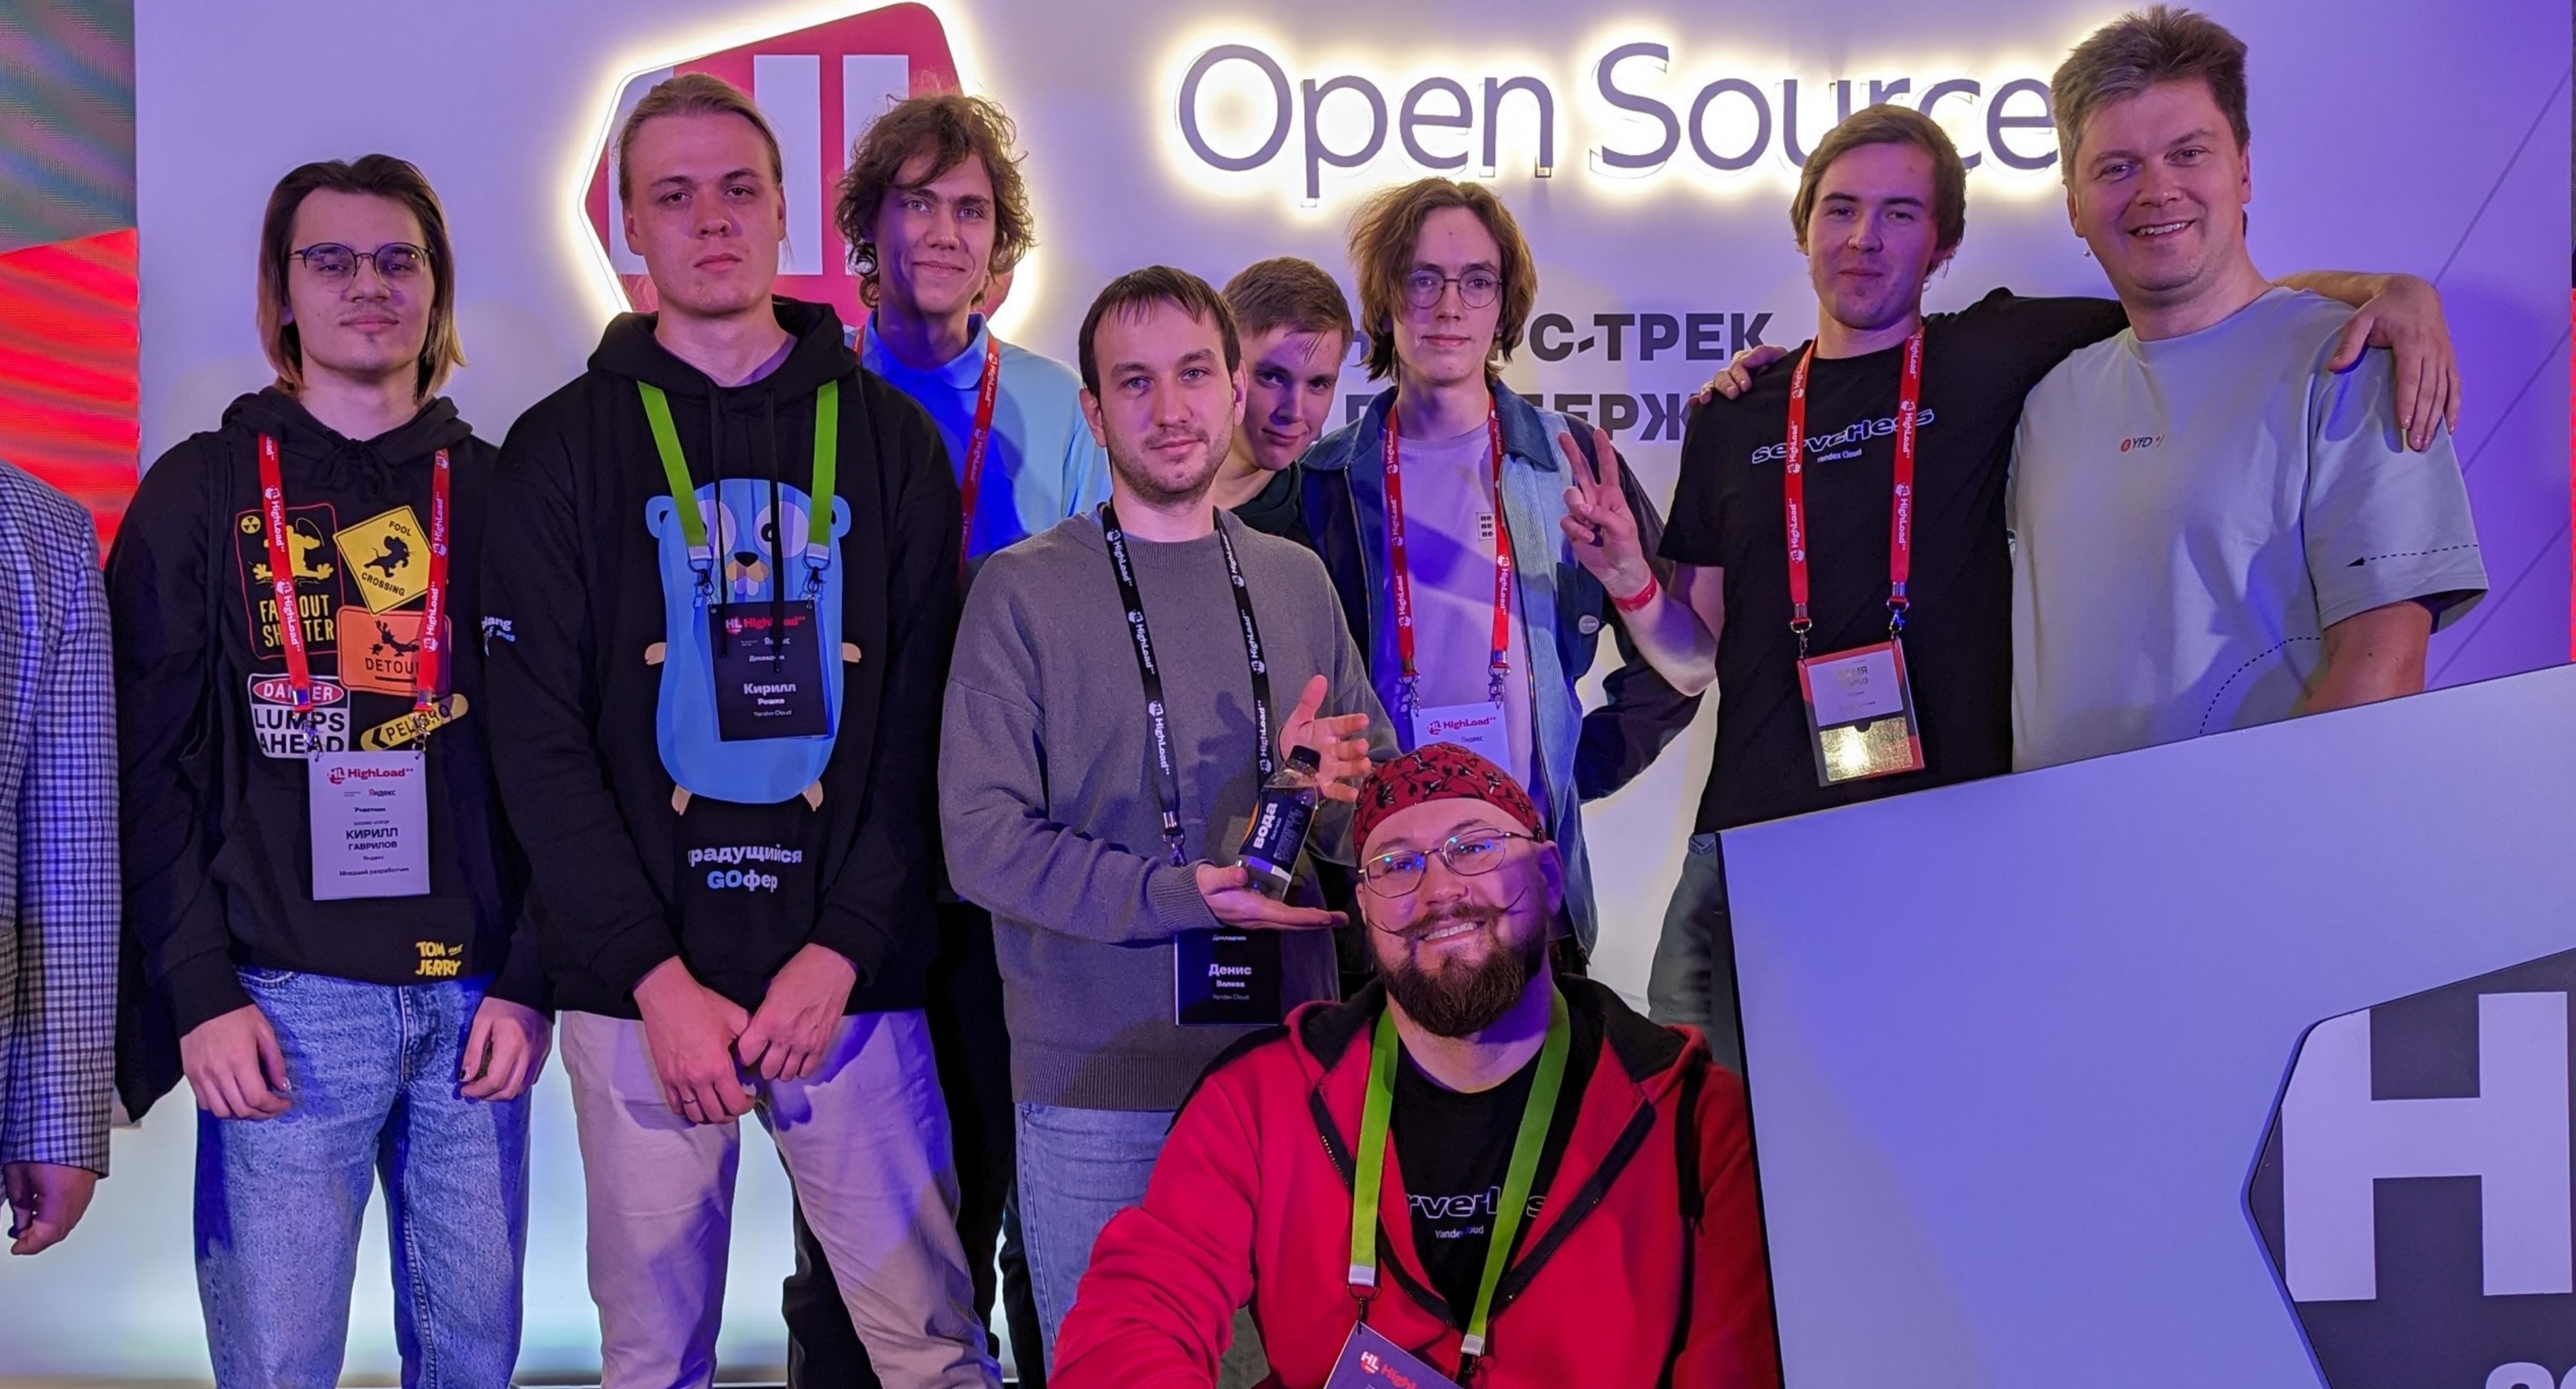


Разрабатывал Postgres, GreenplumDB, Redshift

Мэйнтейнер:

- › Odyssey
- › WAL-G
- › SPQR

Open Source



Чего хотят разработчики?

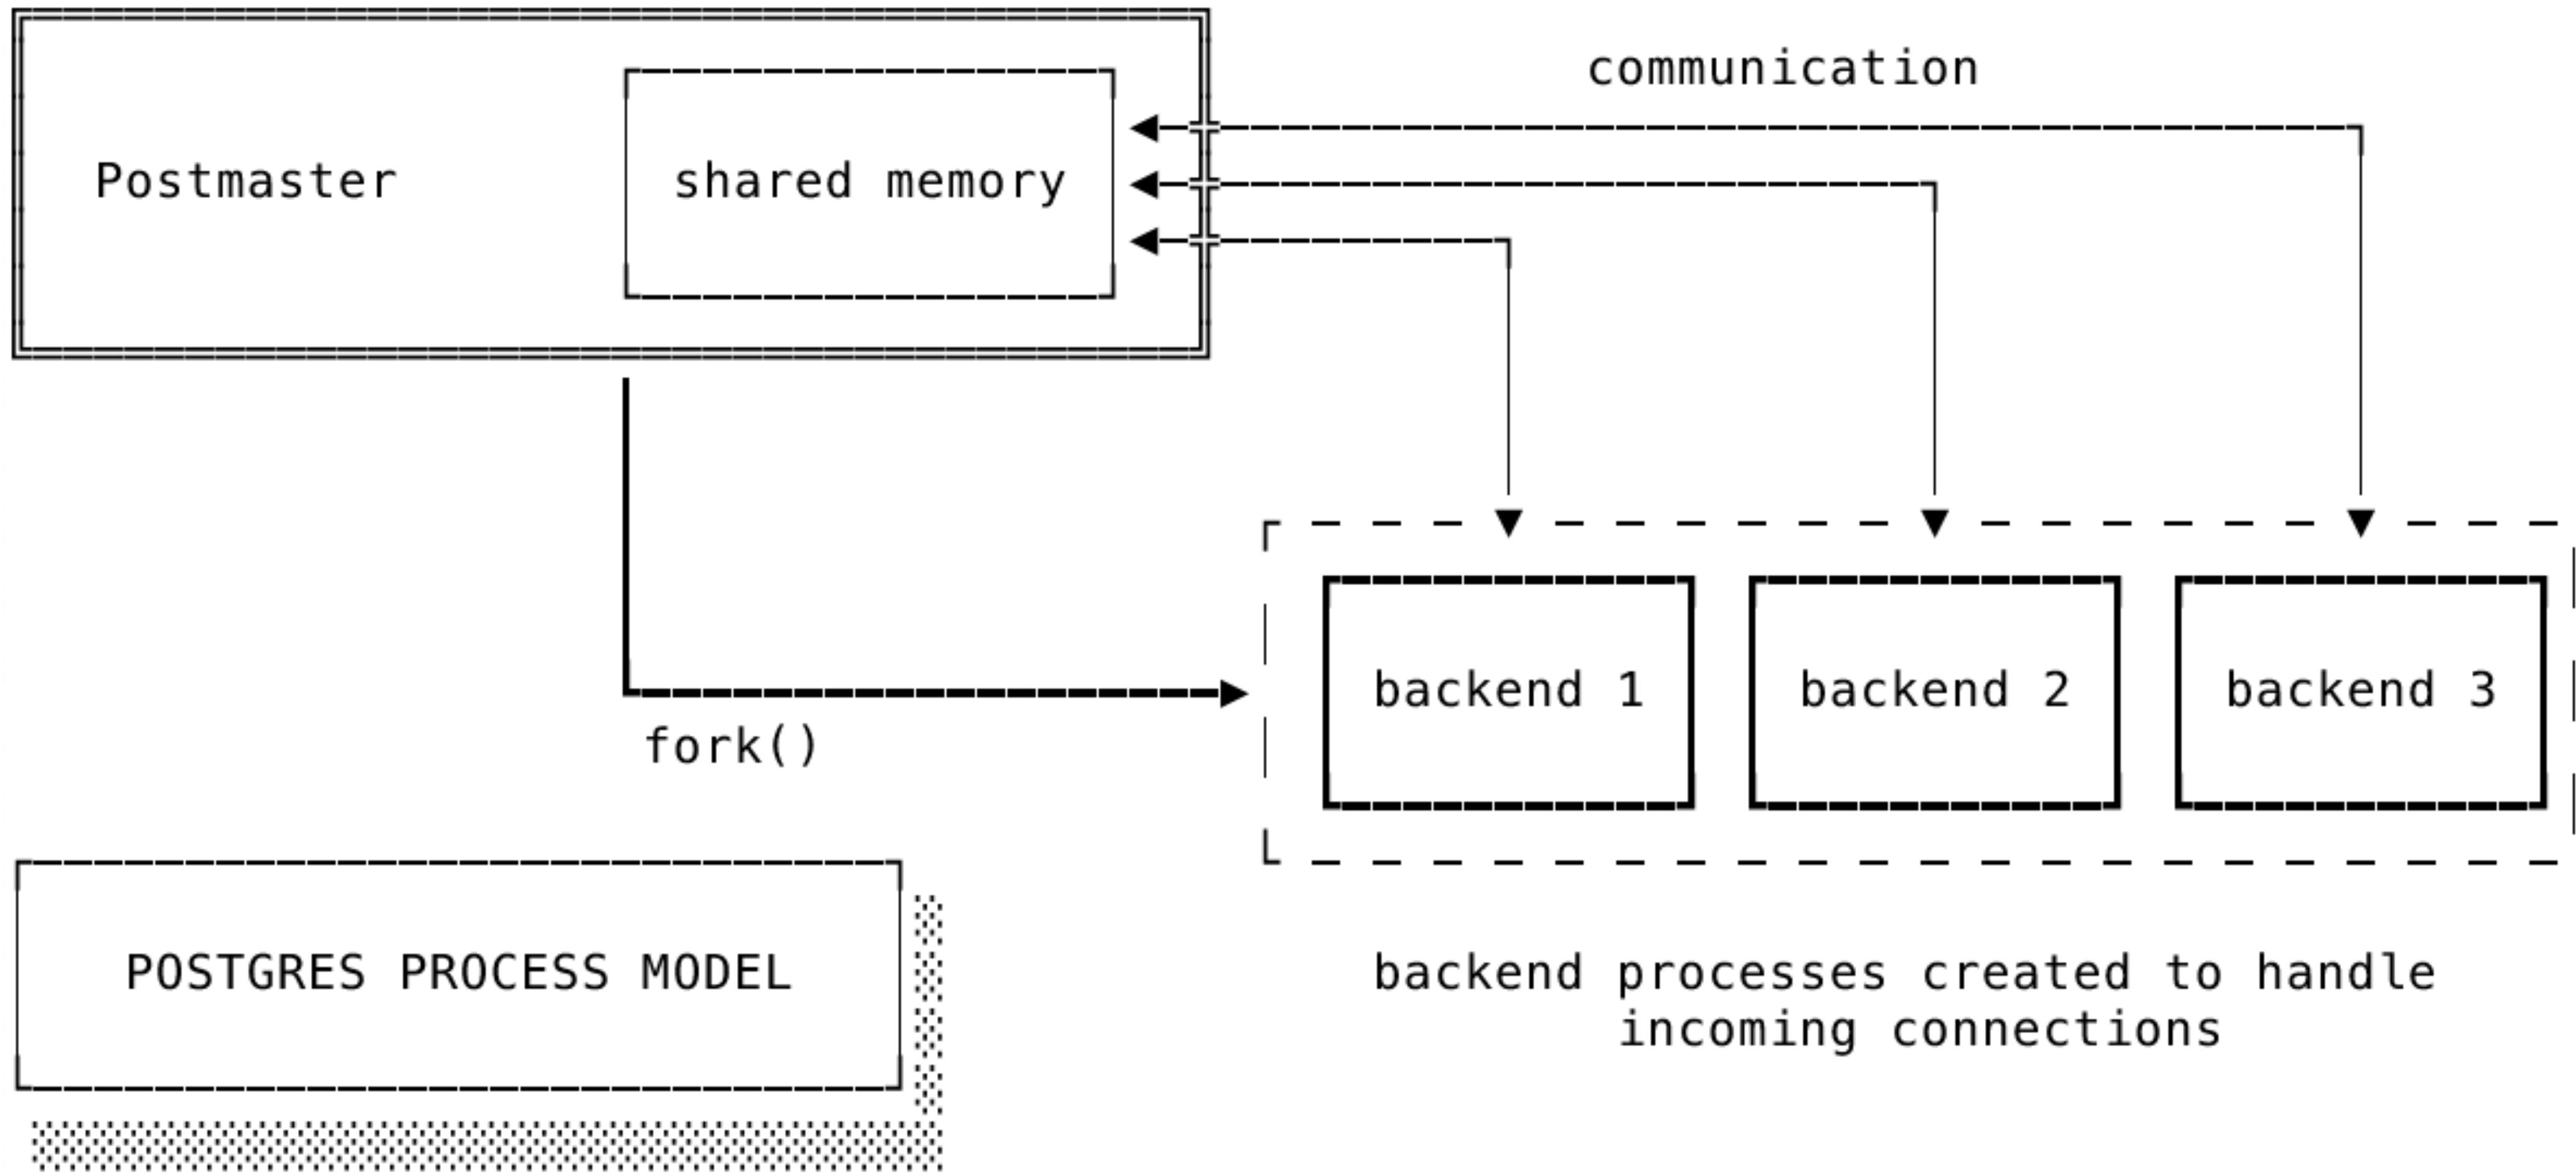
| **Request-response**

Чего хотят разработчики?

| Request-response

А не:

- › Startup cost
- › Dealing with reconnection
- › Dealing with partial state
- › Managing swarm of connections in app



Postgres on the wire

A look at the PostgreSQL wire protocol

Jan Urbański

j.urbanski@wulczer.org

Ducksboard

PGCon 2014, Ottawa, May 23

Protocol versions

- ▶ the 2.0 protocol got introduced in 6.4, around 1999
 - ▶ protocol versioning got added in the previous release
- ▶ the 3.0 got introduced in 7.4, in 2003
- ▶ the server **still supports** protocol 1.0!
- ▶ 3.0 has some new features
 - ▶ extended query protocol
 - ▶ COPY improvements
 - ▶ overall better frame structure

FEBE frame format

Virtually all messages start with an ASCII identifier, followed by length and payload.

Regular packet

| | | |
|----------|-----------|---------|
| char tag | int32 len | payload |
|----------|-----------|---------|

The exception is the startup packet, which starts with the length followed by the protocol version.

Startup packet

| | | |
|-----------|----------------|---------|
| int32 len | int32 protocol | payload |
|-----------|----------------|---------|

```

75  func (p *PgFortuneBackend) handleStartup() error {
76      startupMessage, err := p.backend.ReceiveStartupMessage()
77      if err != nil {
78          return fmt.Errorf("error receiving startup message: %w", err)
79      }
80
81      switch startupMessage.(type) {
82      case *pgproto3.StartupMessage:
83          buf := mustEncode((&pgproto3.AuthenticationOk{}).Encode(nil))
84          buf = mustEncode((&pgproto3.ReadyForQuery{TxStatus: 'I'}).Encode(buf))
85          _, err = p.conn.Write(buf)
86          if err != nil {
87              return fmt.Errorf("error sending ready for query: %w", err)
88          }
89      case *pgproto3.SSLRequest:
90          _, err = p.conn.Write([]byte("N"))
91          if err != nil {
92              return fmt.Errorf("error sending deny SSL request: %w", err)
93          }
94          return p.handleStartup()
95      default:
96          return fmt.Errorf("unknown startup message: %#v", startupMessage)
97      }
98
99      return nil
100 }

```

Cancellation

Cancel request

| | | | |
|-----------|------------------|-----------|--------------|
| int32 len | int32 cancelcode | int32 pid | int32 secret |
|-----------|------------------|-----------|--------------|

- ▶ the **cancel key** is transmitted by the server upon connection
- ▶ cancelling queries requires opening **separate connection**
- ▶ another dummy protocol version is sent to ask for cancellation
- ▶ the cancellation message includes the process ID and a 32 bit key
 - ▶ theoretically open to **replay attacks**, but can be sent over SSL
 - ▶ libpq **does not**, so most applications will transmit it in the open

Handling errors

ErrorResponse

| | | | | | | | | | |
|-----|-----------|-----------|-----------|----|-----------|-----------|----|-----|----|
| 'E' | int32 len | char code | str value | \0 | char code | str value | \0 | ... | \0 |
|-----|-----------|-----------|-----------|----|-----------|-----------|----|-----|----|

- ▶ the ErrorResponse message is sent for all kinds of errors
 - ▶ both for authentication errors and client errors
- ▶ it is a list of key-value fields
 - ▶ in 2.0 it was just a string, in 3.0 it has structure
 - ▶ example error fields are: message, detail, hint and error position
 - ▶ detailed down to the source file and line, great for fingerprinting

Simple query frames

Query

| | | |
|-----|-----------|-----------|
| 'Q' | int32 len | str query |
|-----|-----------|-----------|

RowDescription

| | | | | | | |
|---------|----------------|-----------------|---------------|---------------|--------------|--------------|
| 'T' | int32 len | int16 numfields | + | | | |
| str col | int32 tableoid | int16 colno | int32 typeoid | int16 typelen | int32 typmod | int16 format |

...

DataRow

| | | | | | |
|-----|-----------|-----------------|----------------|---------------------|-----|
| 'D' | int32 len | int16 numfields | int32 fieldlen | char[fieldlen] data | ... |
|-----|-----------|-----------------|----------------|---------------------|-----|

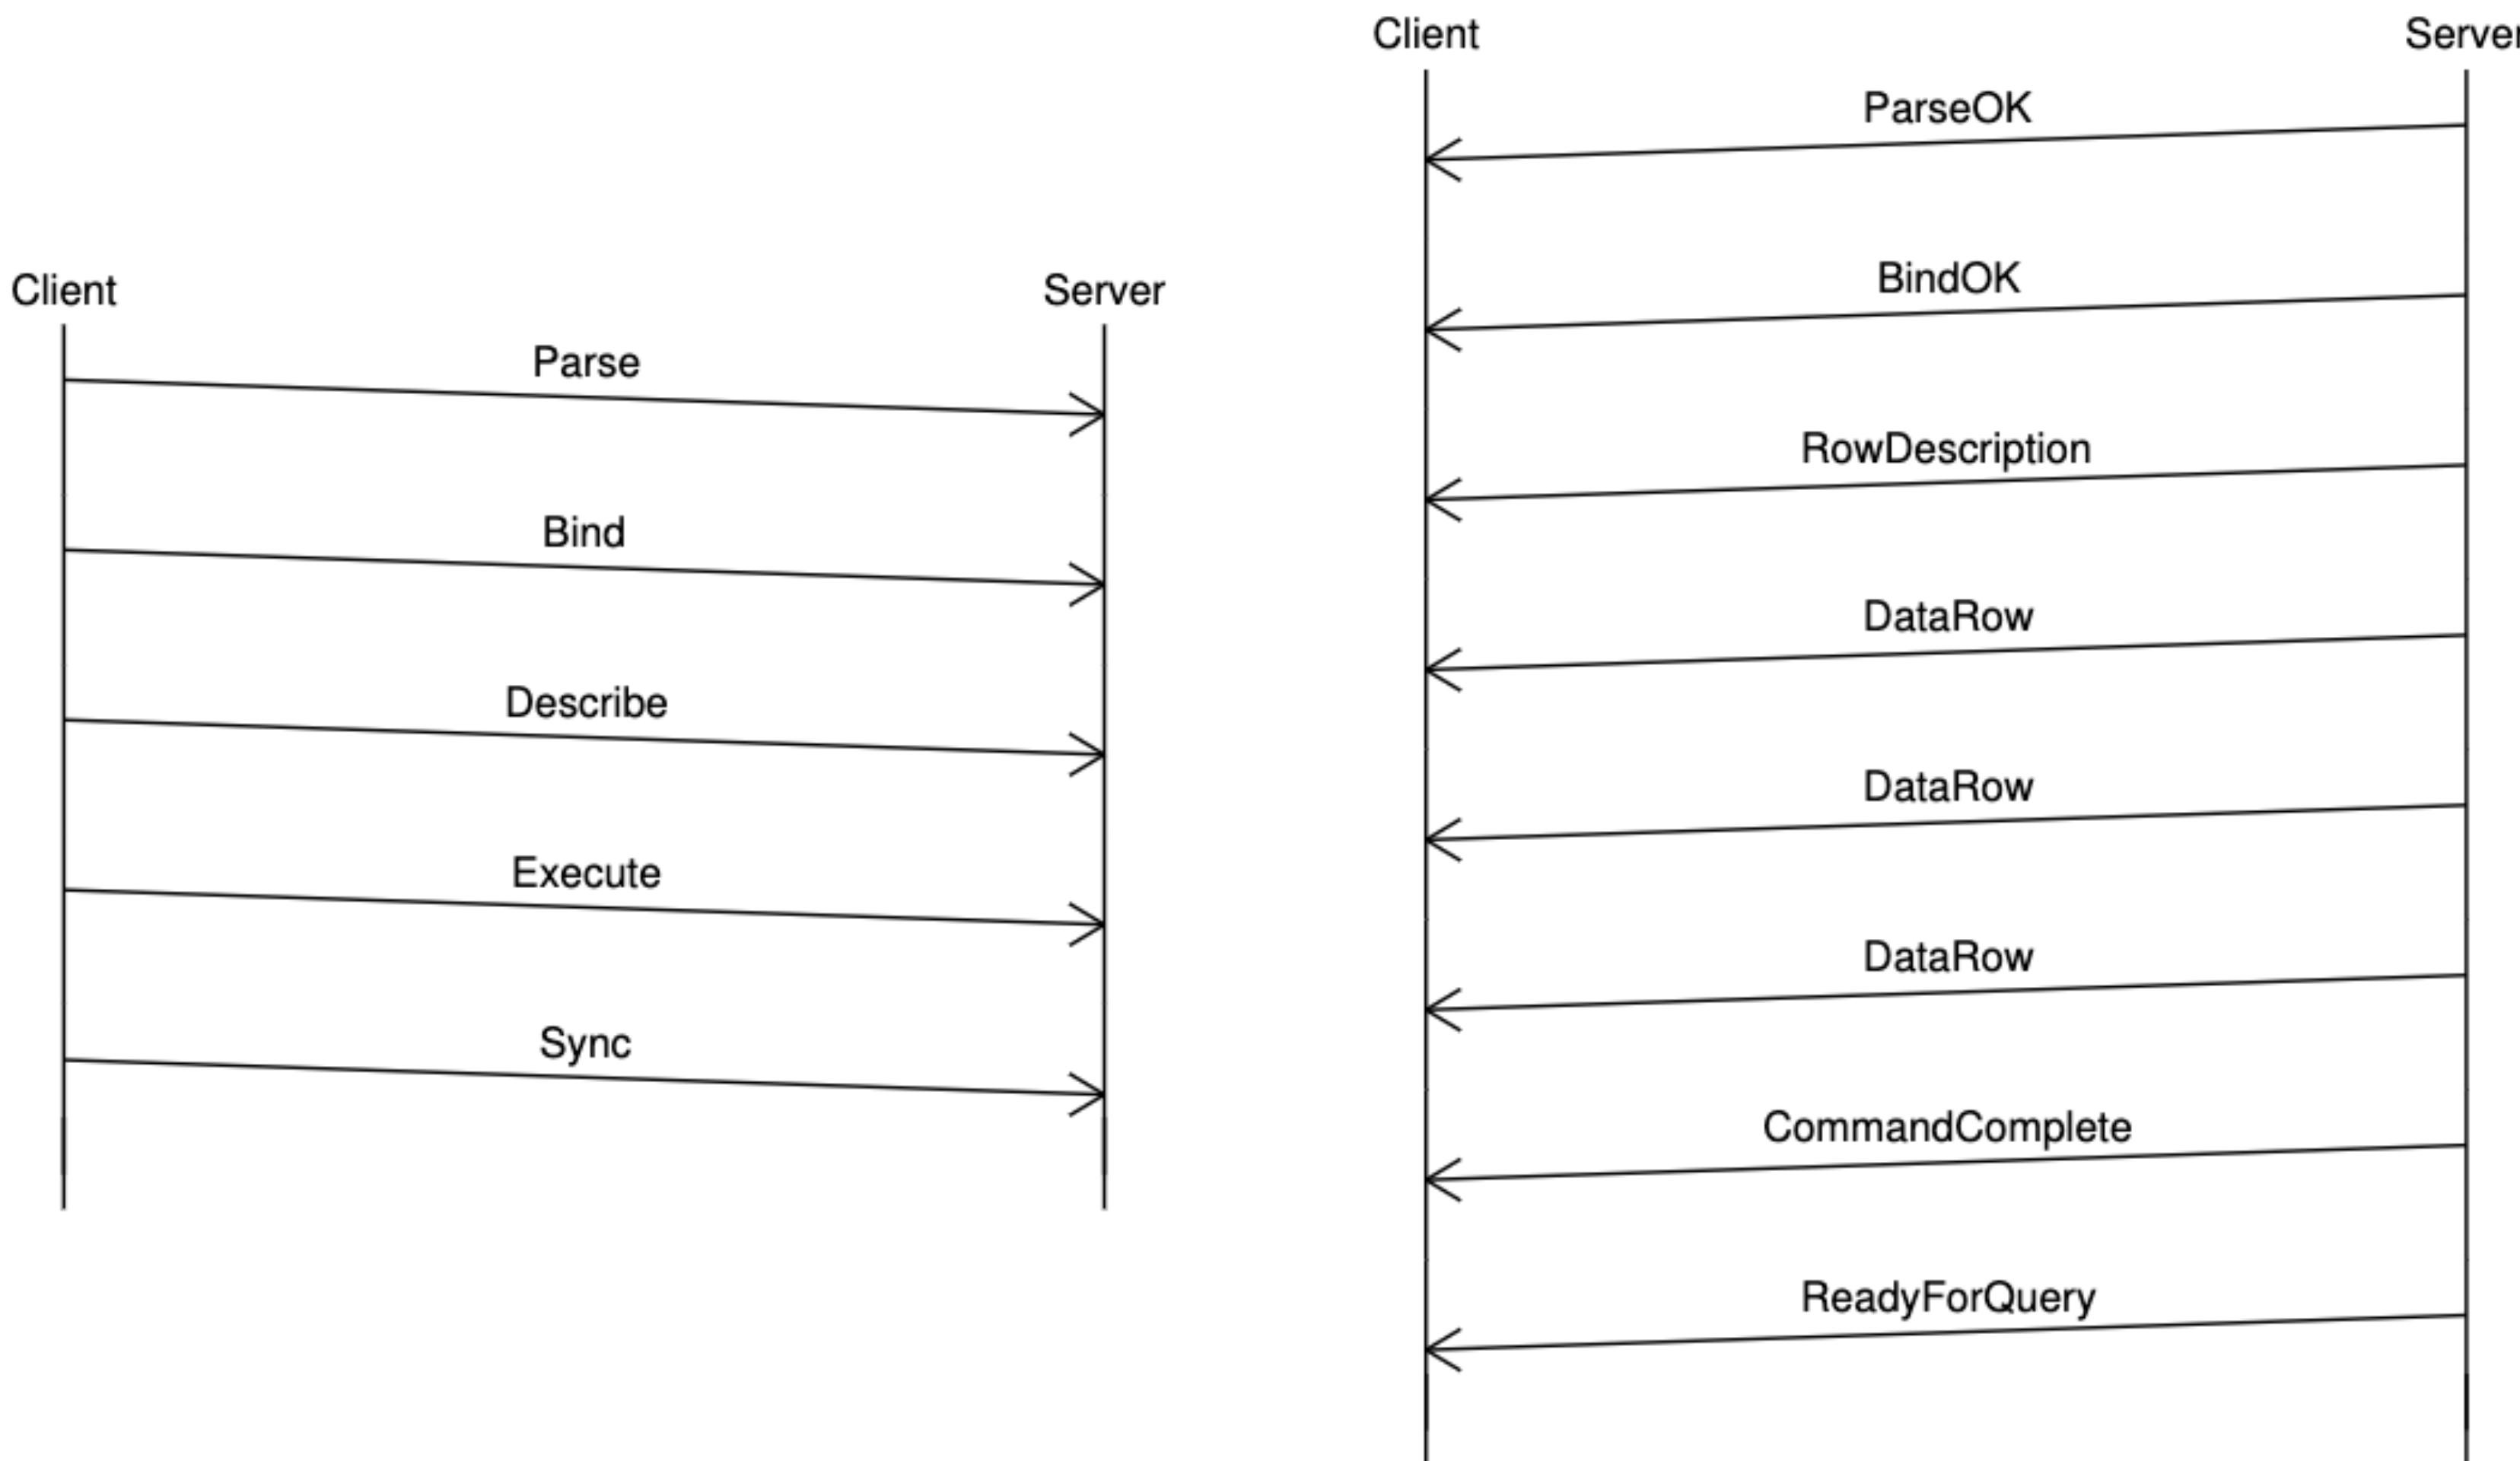
CommandComplete

| | | |
|-----|-----------|---------|
| 'C' | int32 len | str tag |
|-----|-----------|---------|

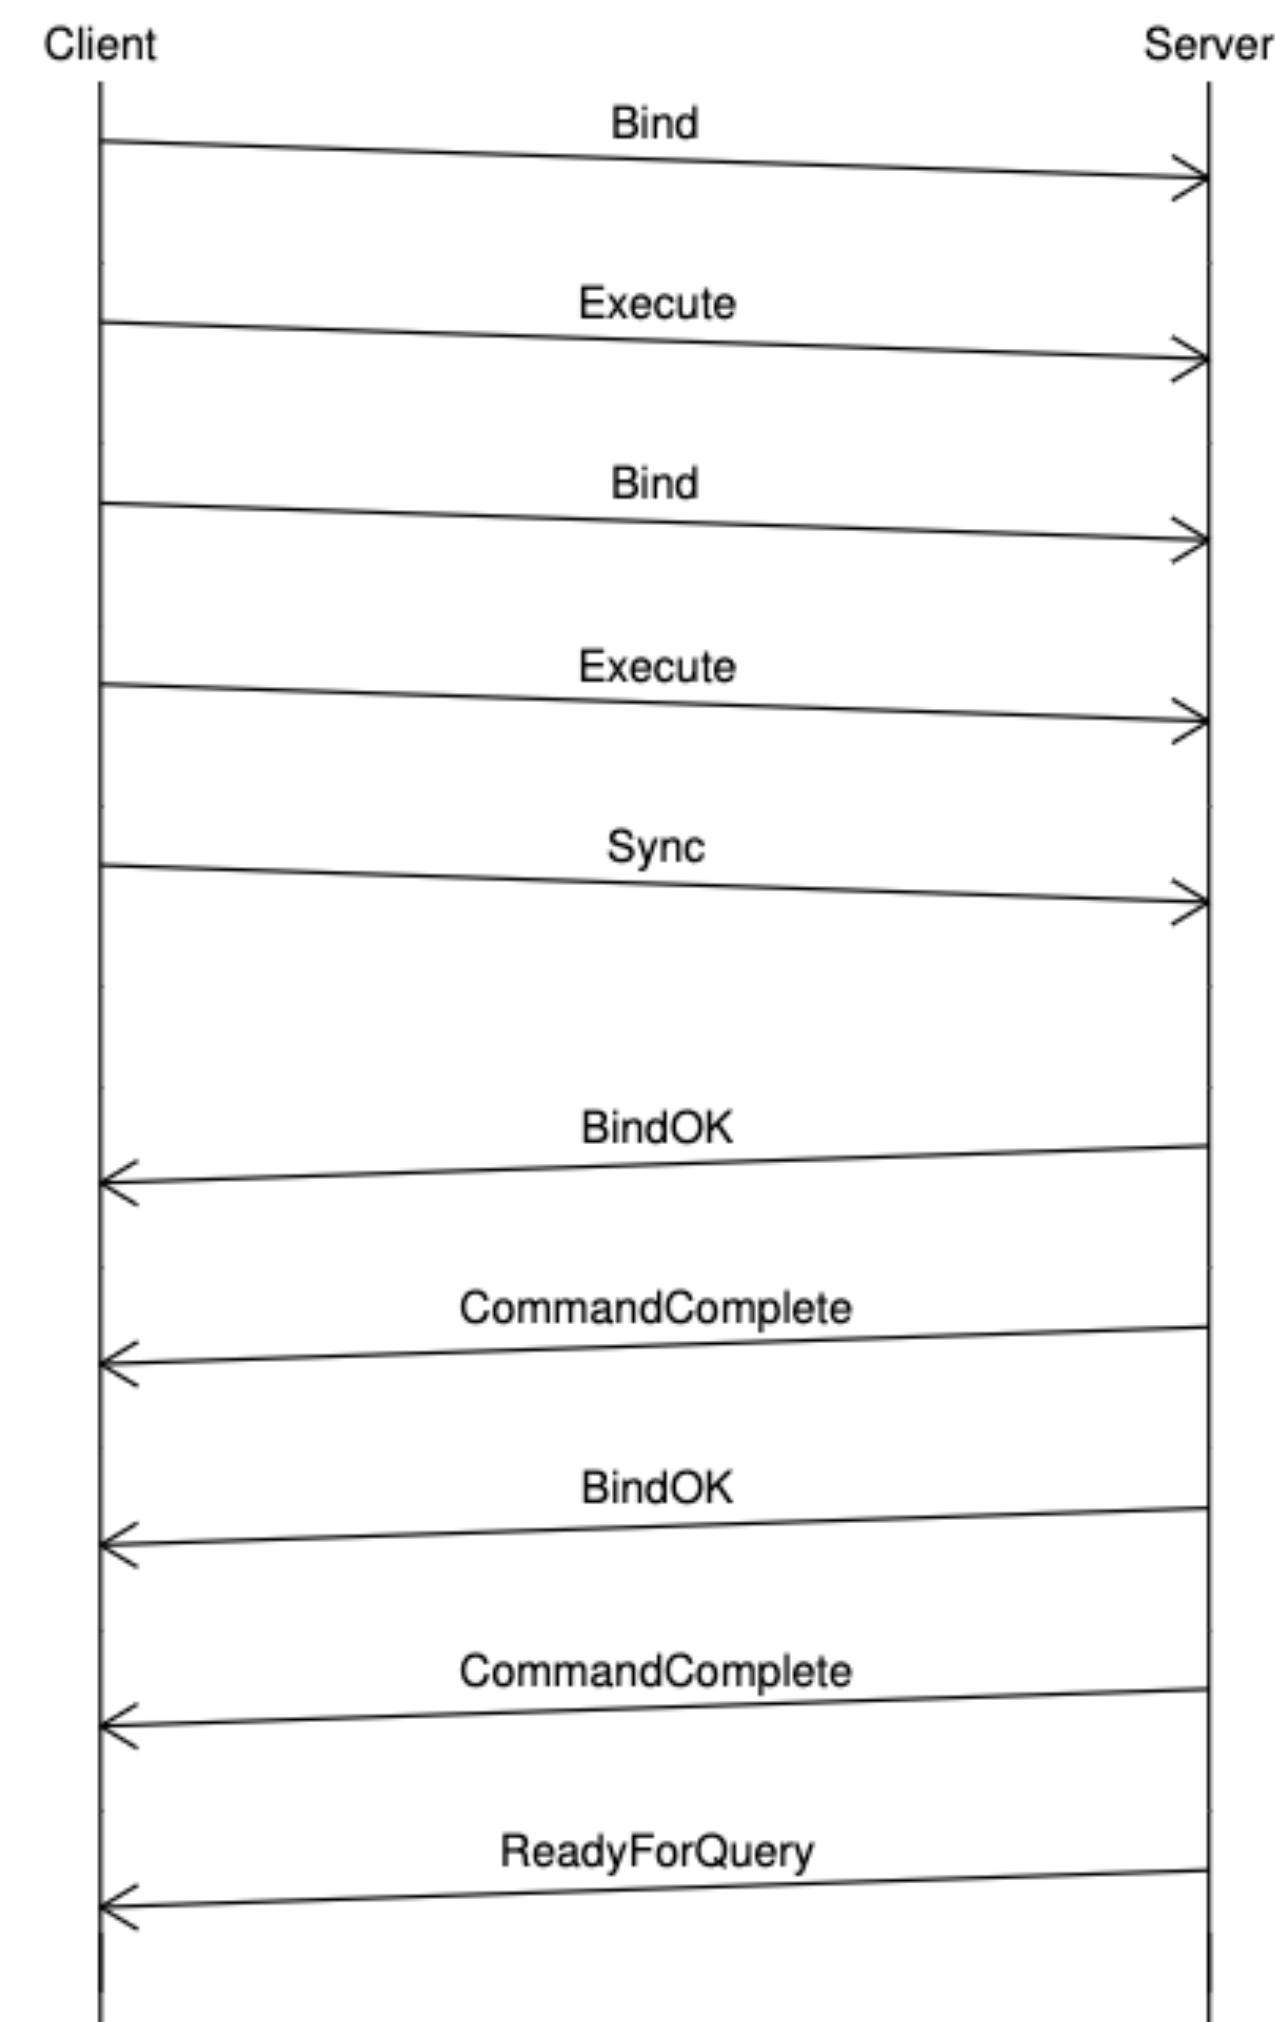
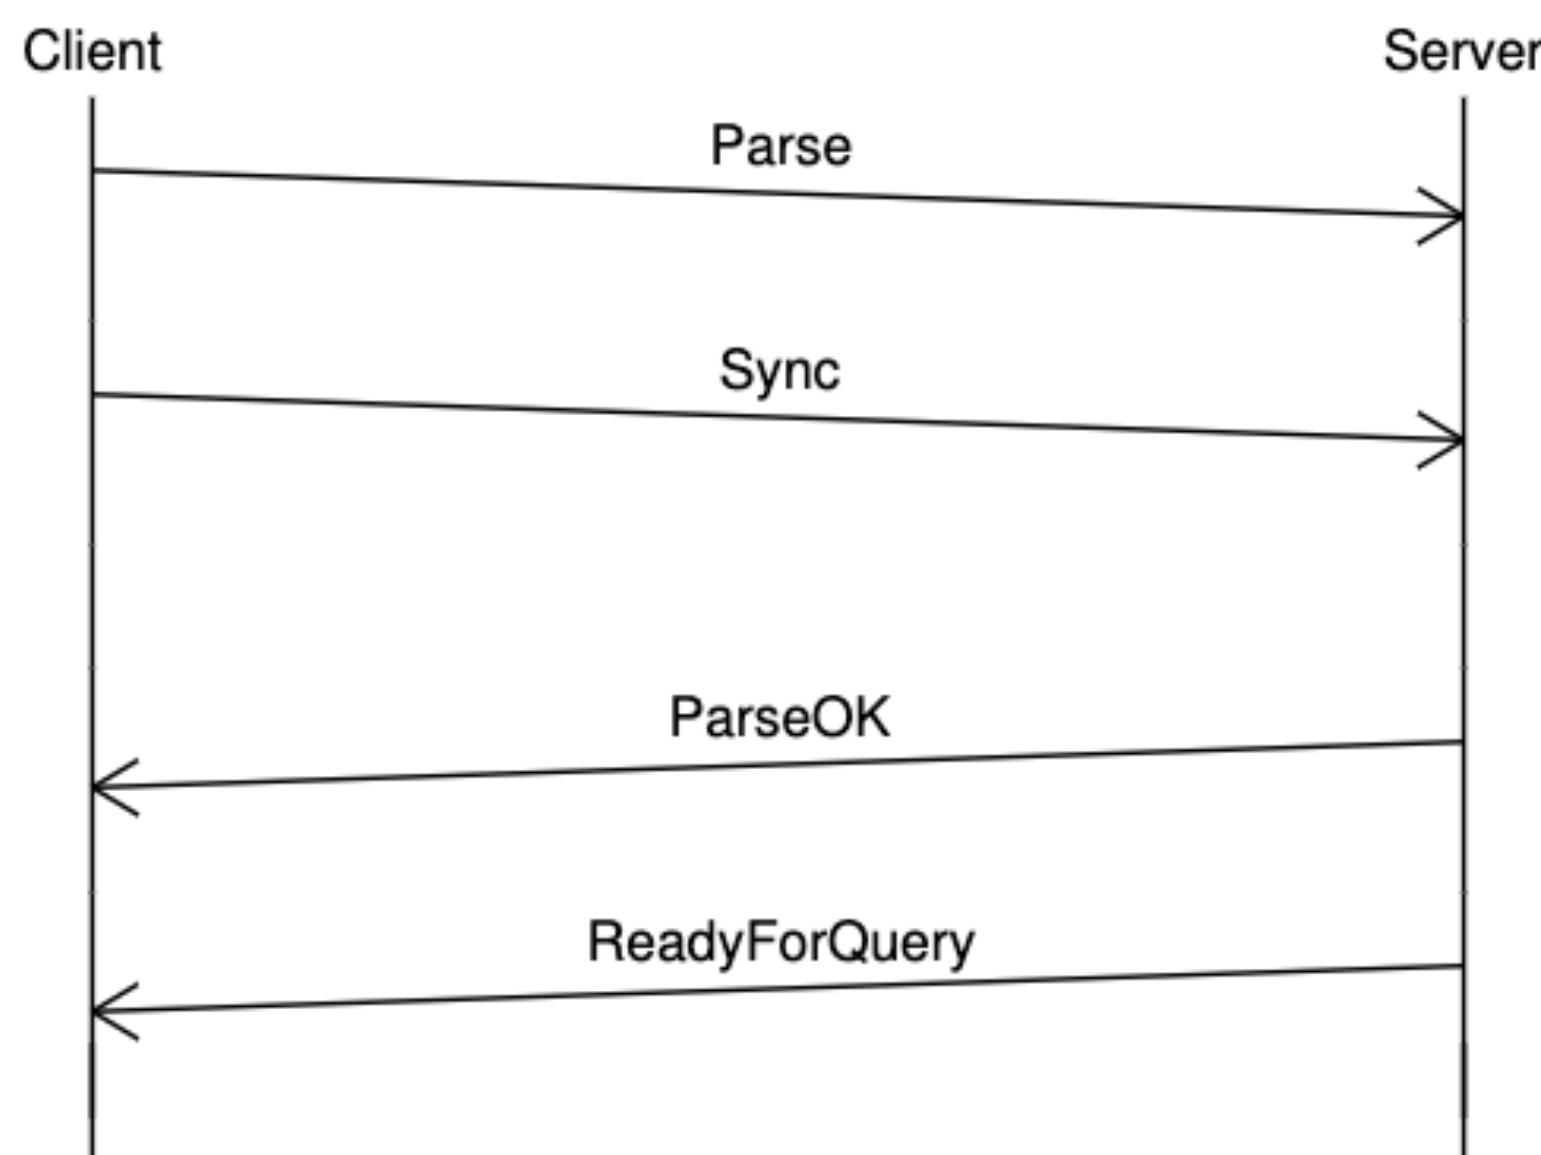
ReadyForQuery

| | | |
|-----|-----------|-------------------|
| 'Z' | int32 len | 'I' or 'T' or 'E' |
|-----|-----------|-------------------|

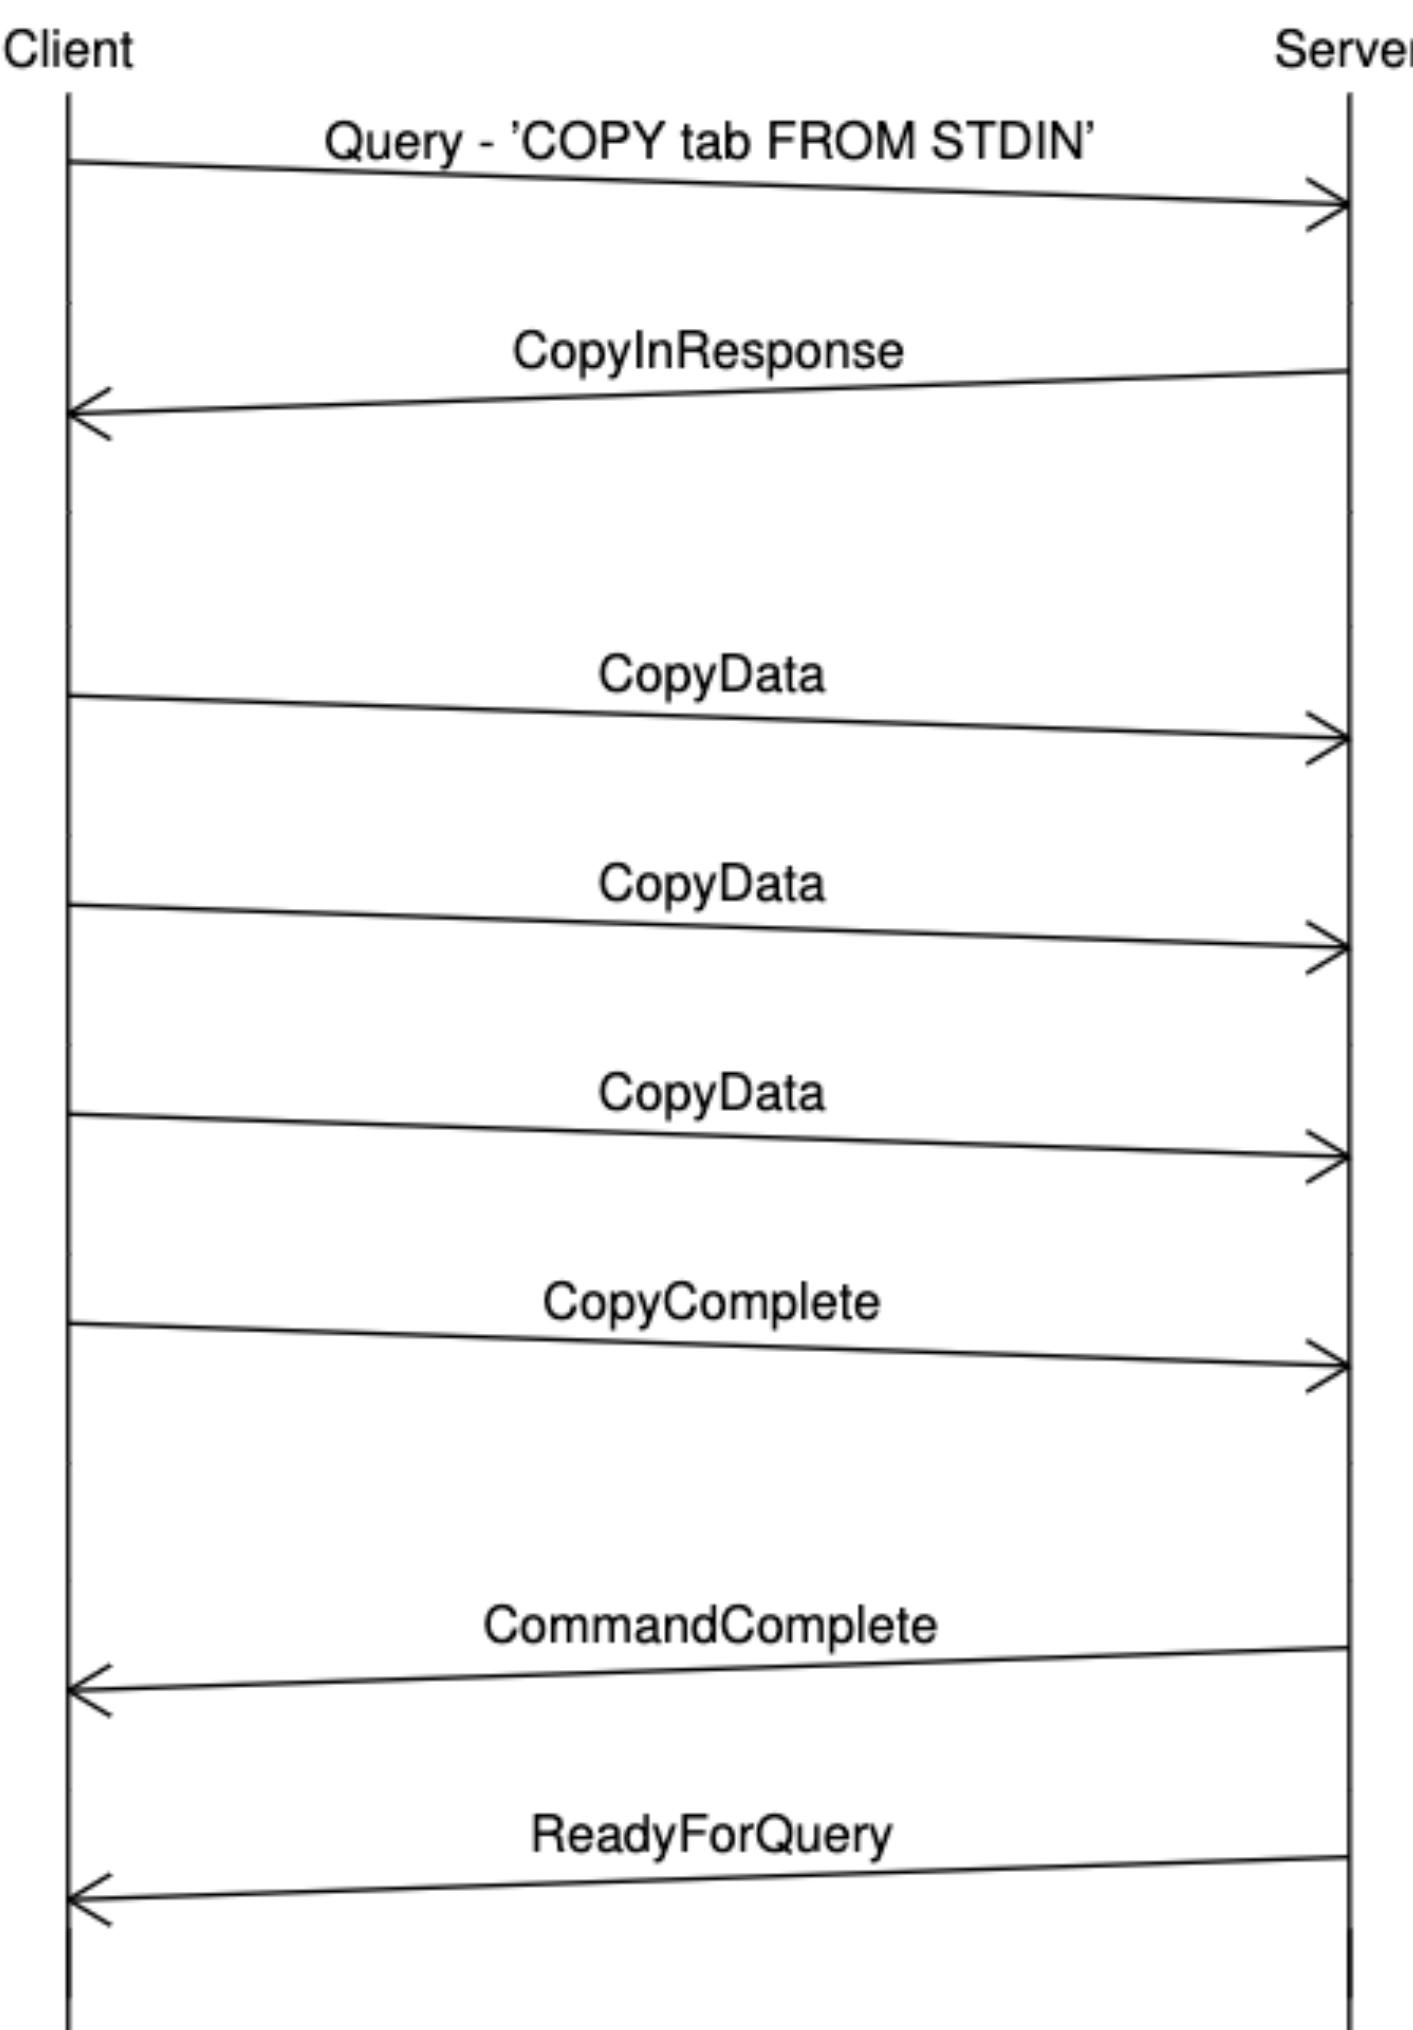
Extended query protocol flow



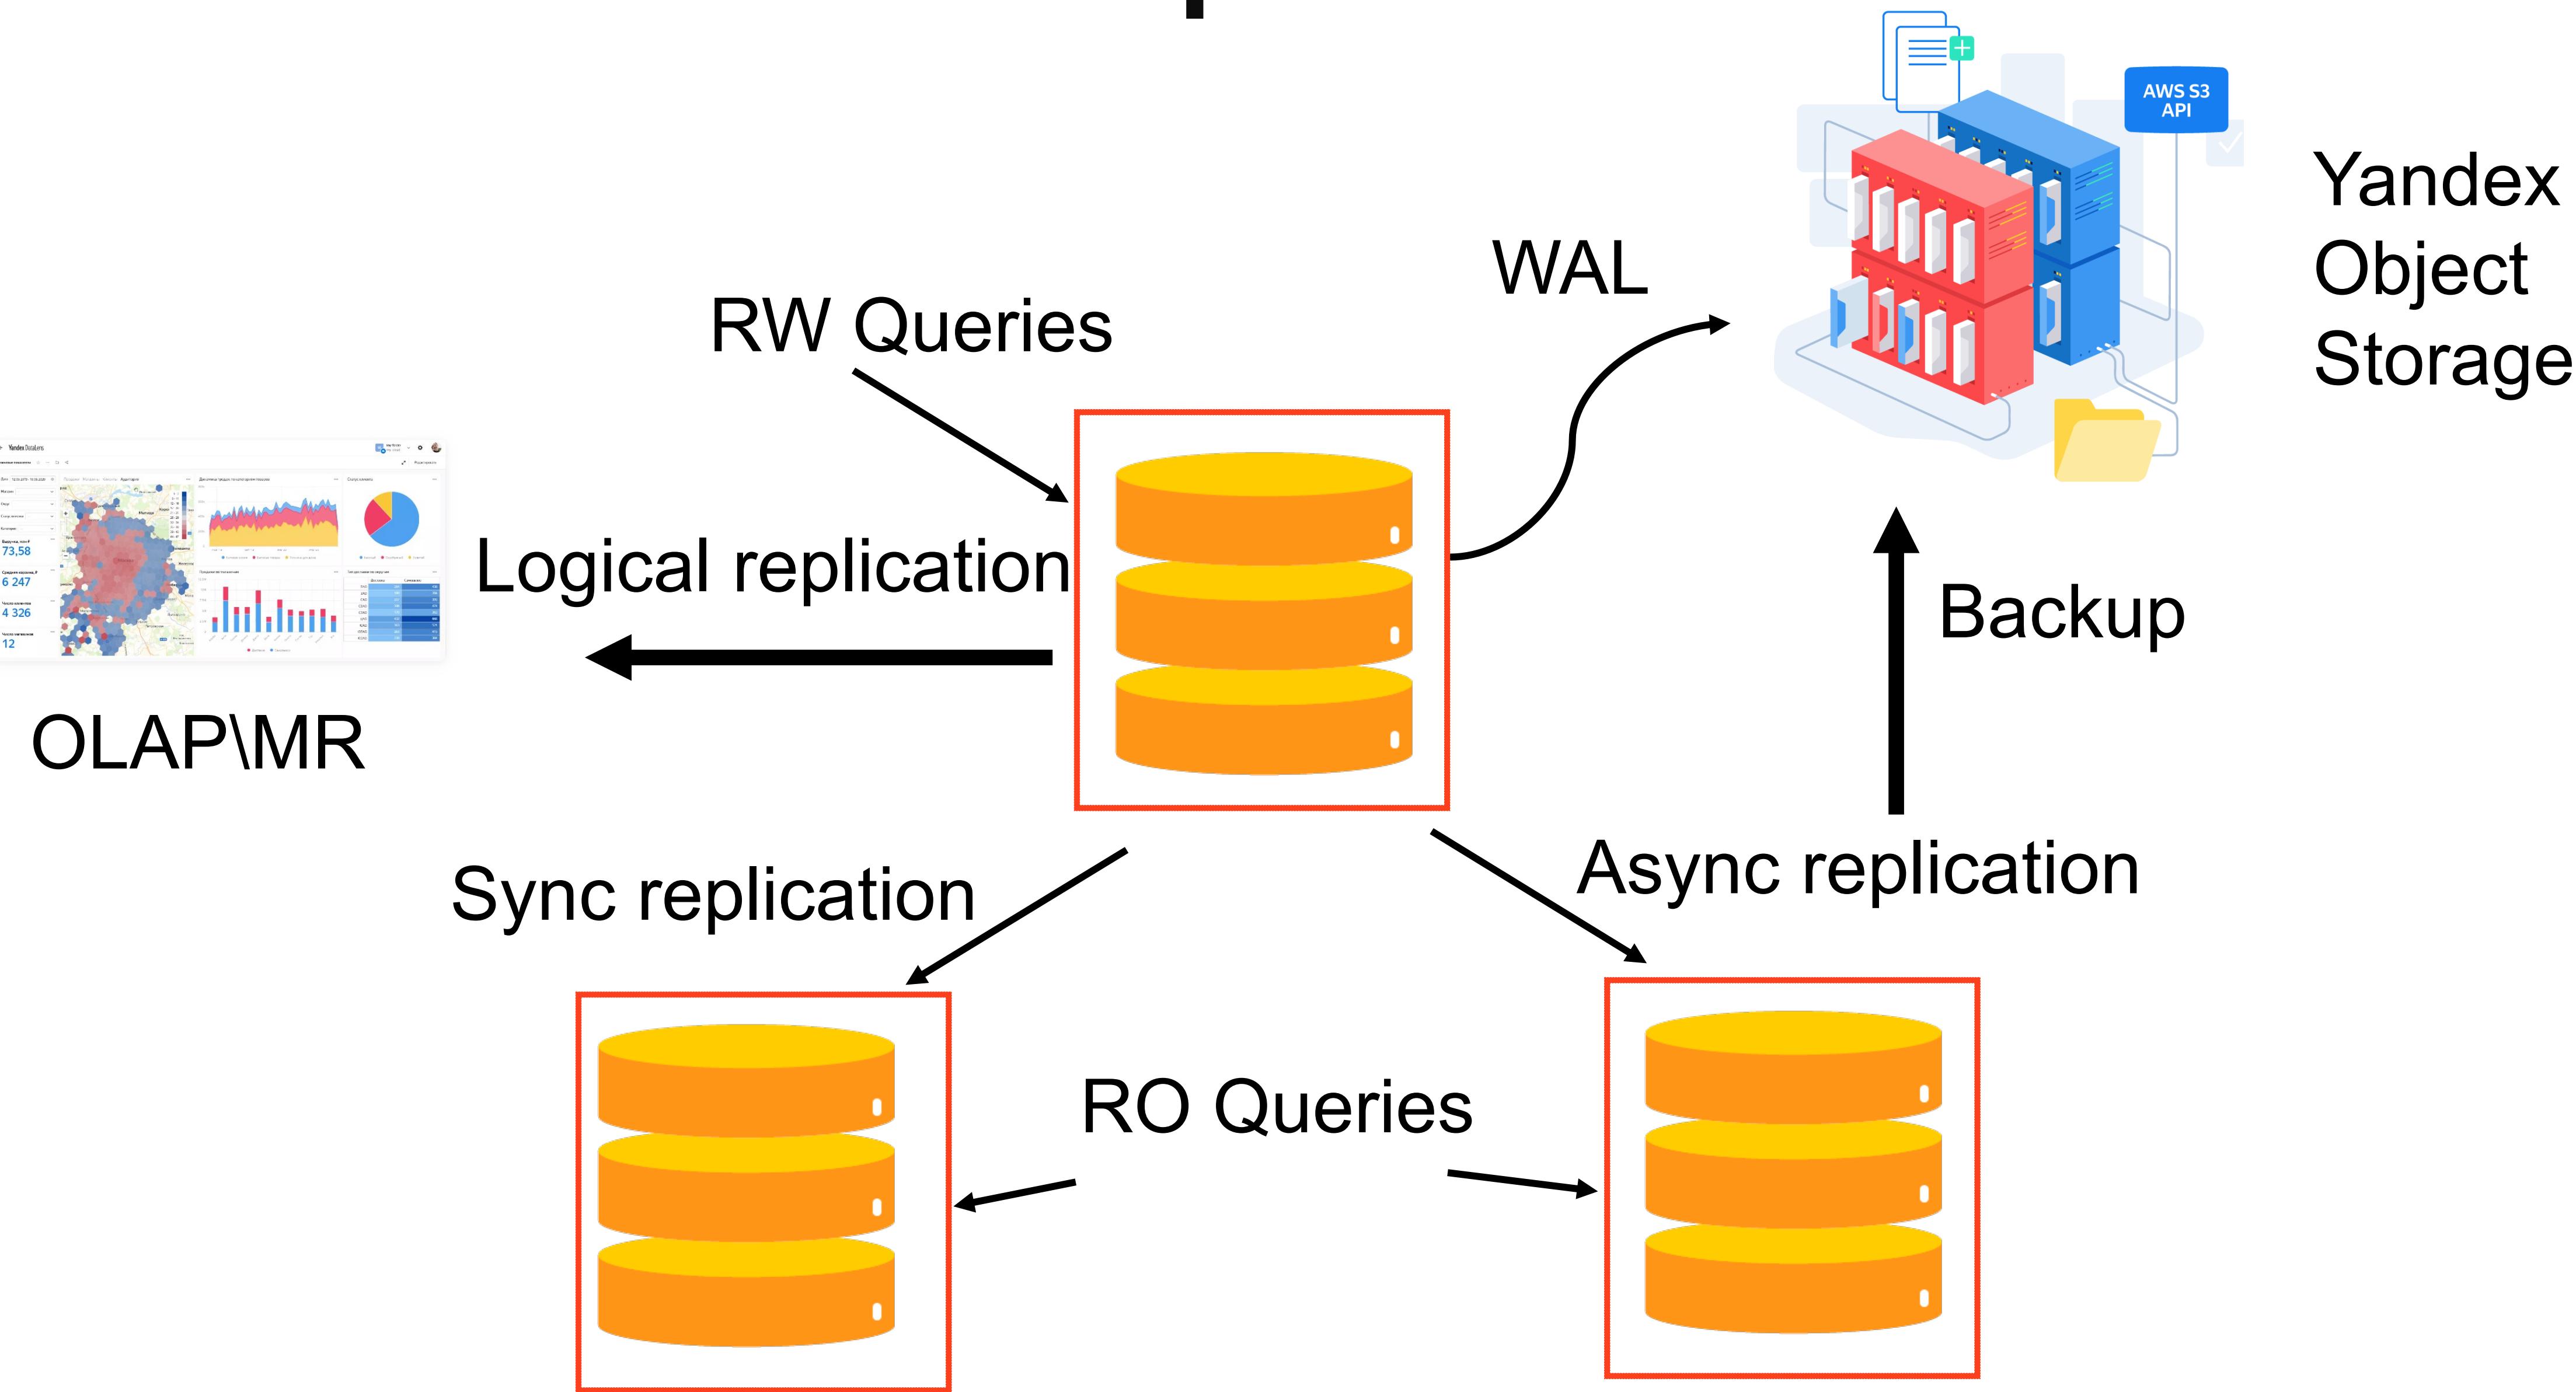
Advanced extended protocol usage



COPY subprotocol flow



Топология кластера в Облаке



Что делает PQ?

- › Запрос-ответ
 - | **text\binary formats**

Что делает PQ?

- › Запрос-ответ
- › FDW
- › Notices

Что делает PQ?

- › Запрос-ответ
- › FDW
- › Notices
- › Listen-Notify

Что делает PQ?

- › Запрос-ответ
- › FDW
- › Notices
- › Listen-Notify
- › Basebackup

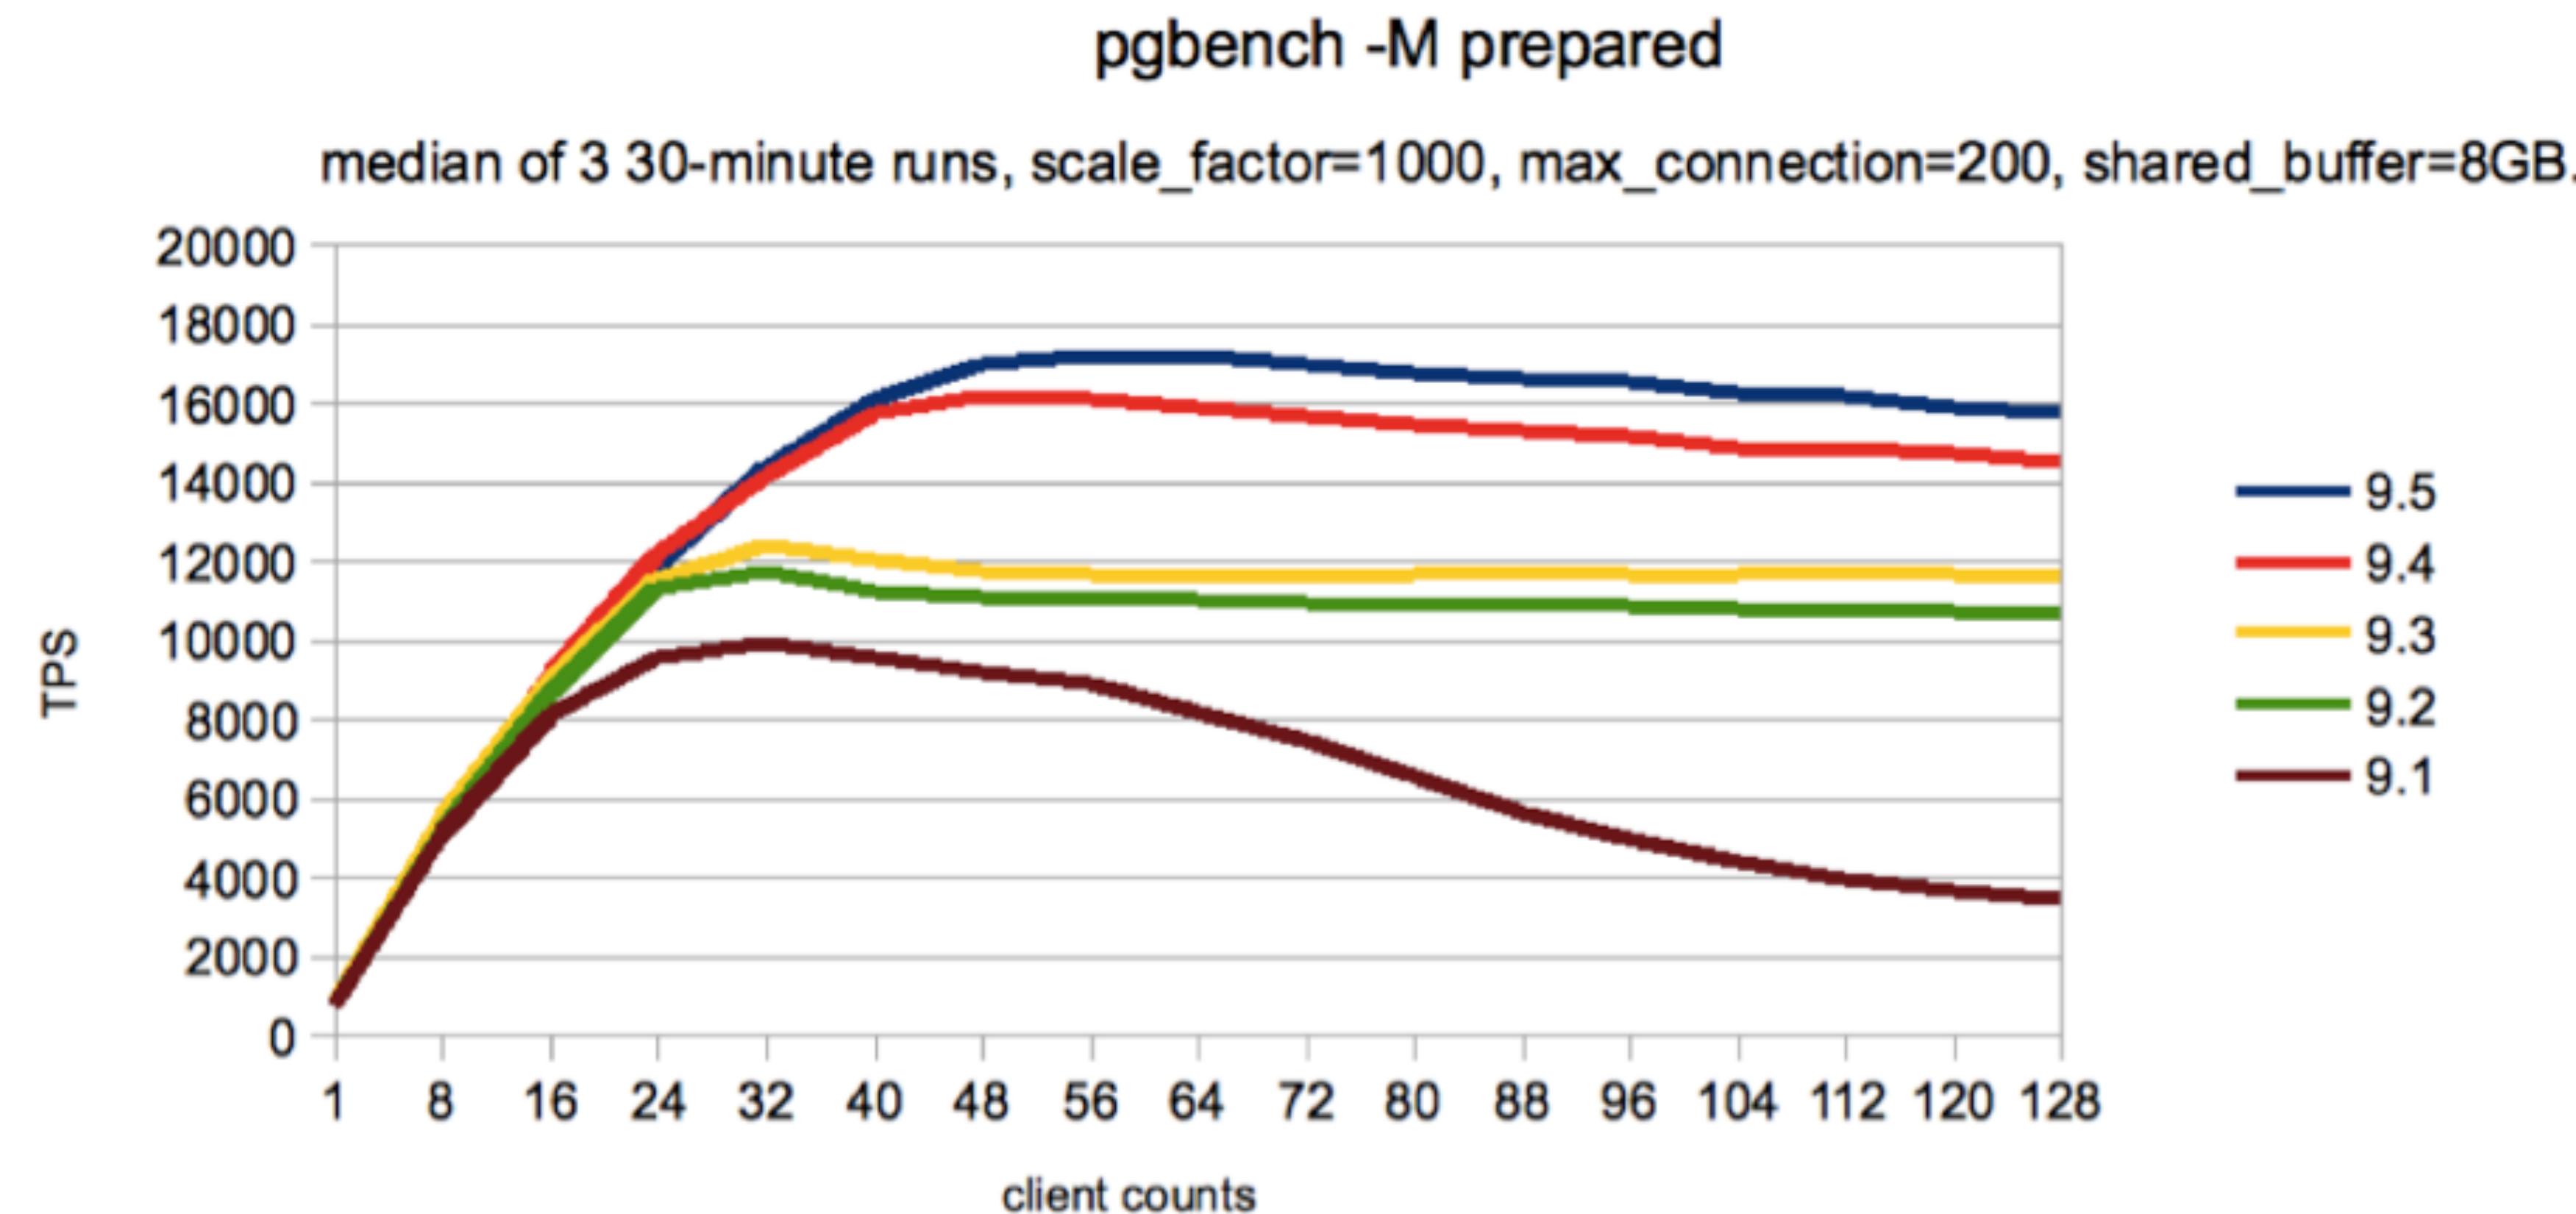
Что делает PQ?

- › Запрос-ответ
- › FDW
- › Notices
- › Listen-Notify
- › Basebackup
- › Walsender\Walreciever

Что делает PQ?

- › Запрос-ответ
- › FDW
- › Notices
- › Listen-Notify
- › Basebackup
- › Walsender\Walreciever
- › Pg_rewind
- › Proxies!

Proxies



Proxies do more!

- › PAUSE + Upgrade
- › Retry
- › Sharding
- › Encryption termination
- etc

Auth methods

- › MD5 auth
- › SCRAM-Auth
- › Cert auth
- › PAM
- › LDAP

CRIME attack: secrets and user data may be compressed together

- Startup (username: website, password: secret_data)
- SimpleQuery("select * from news where user_id = \$1", user_data)

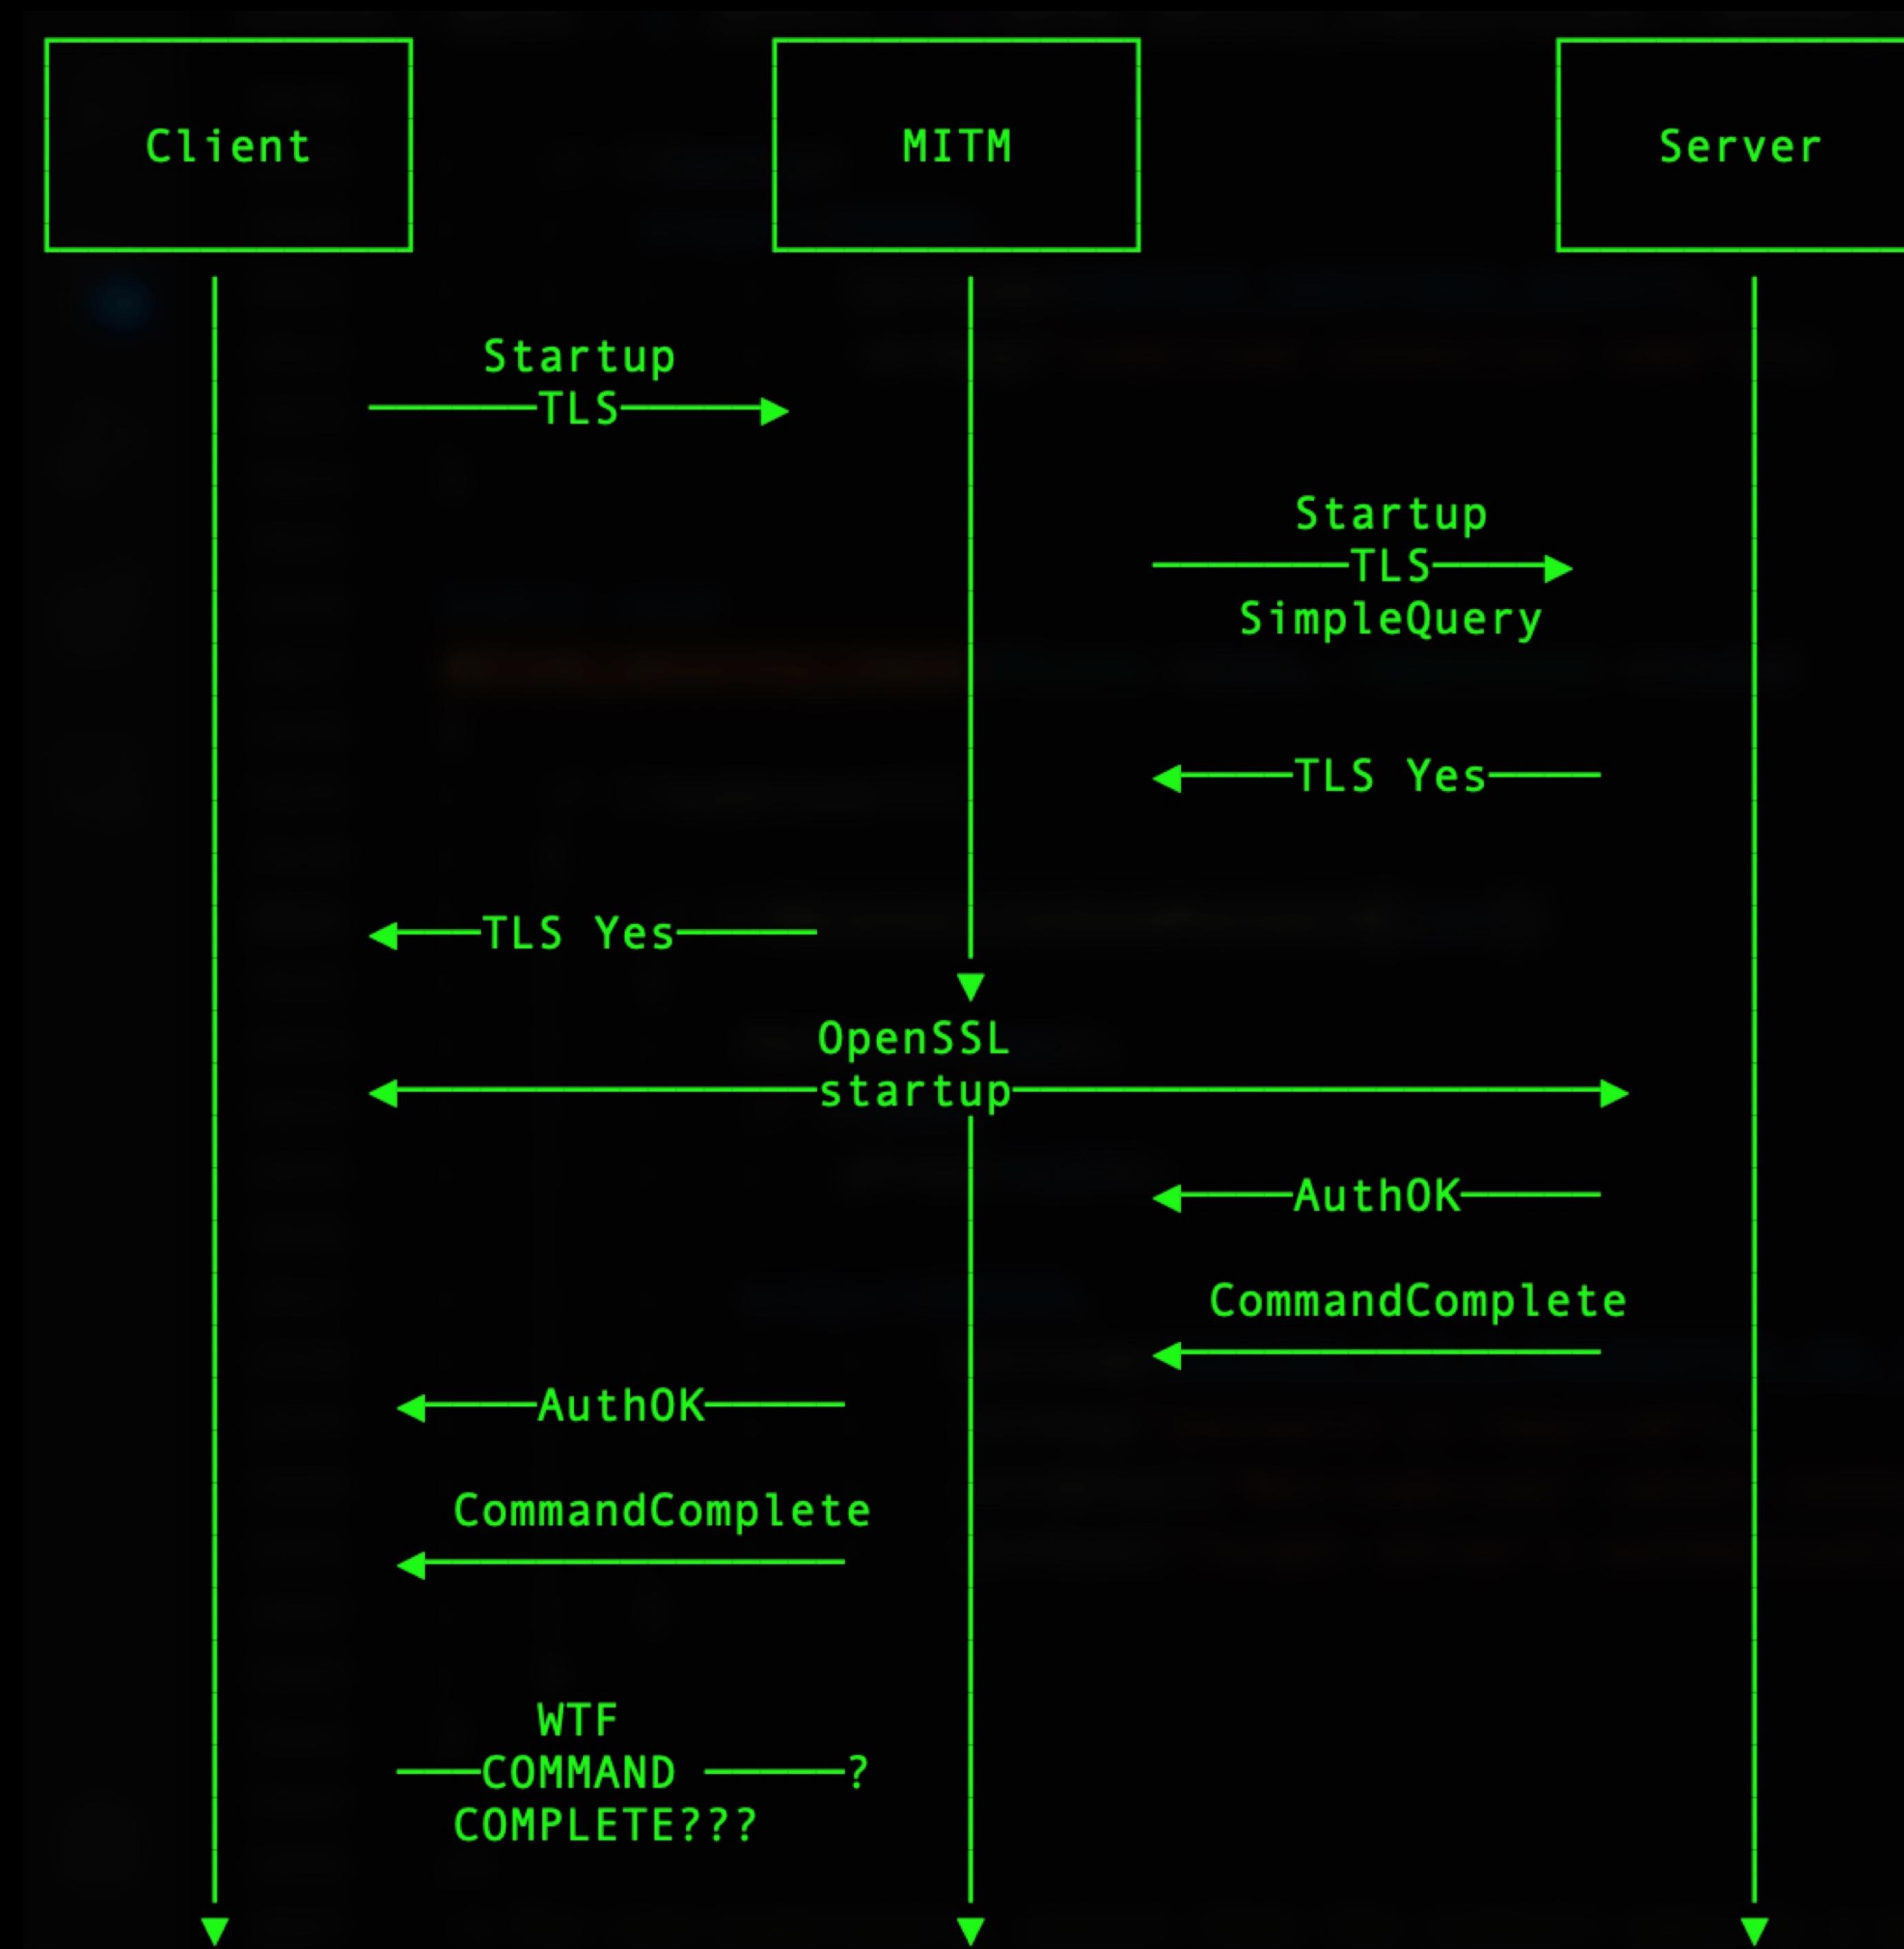
Length of resulting package reveals if user data is similar to secret data

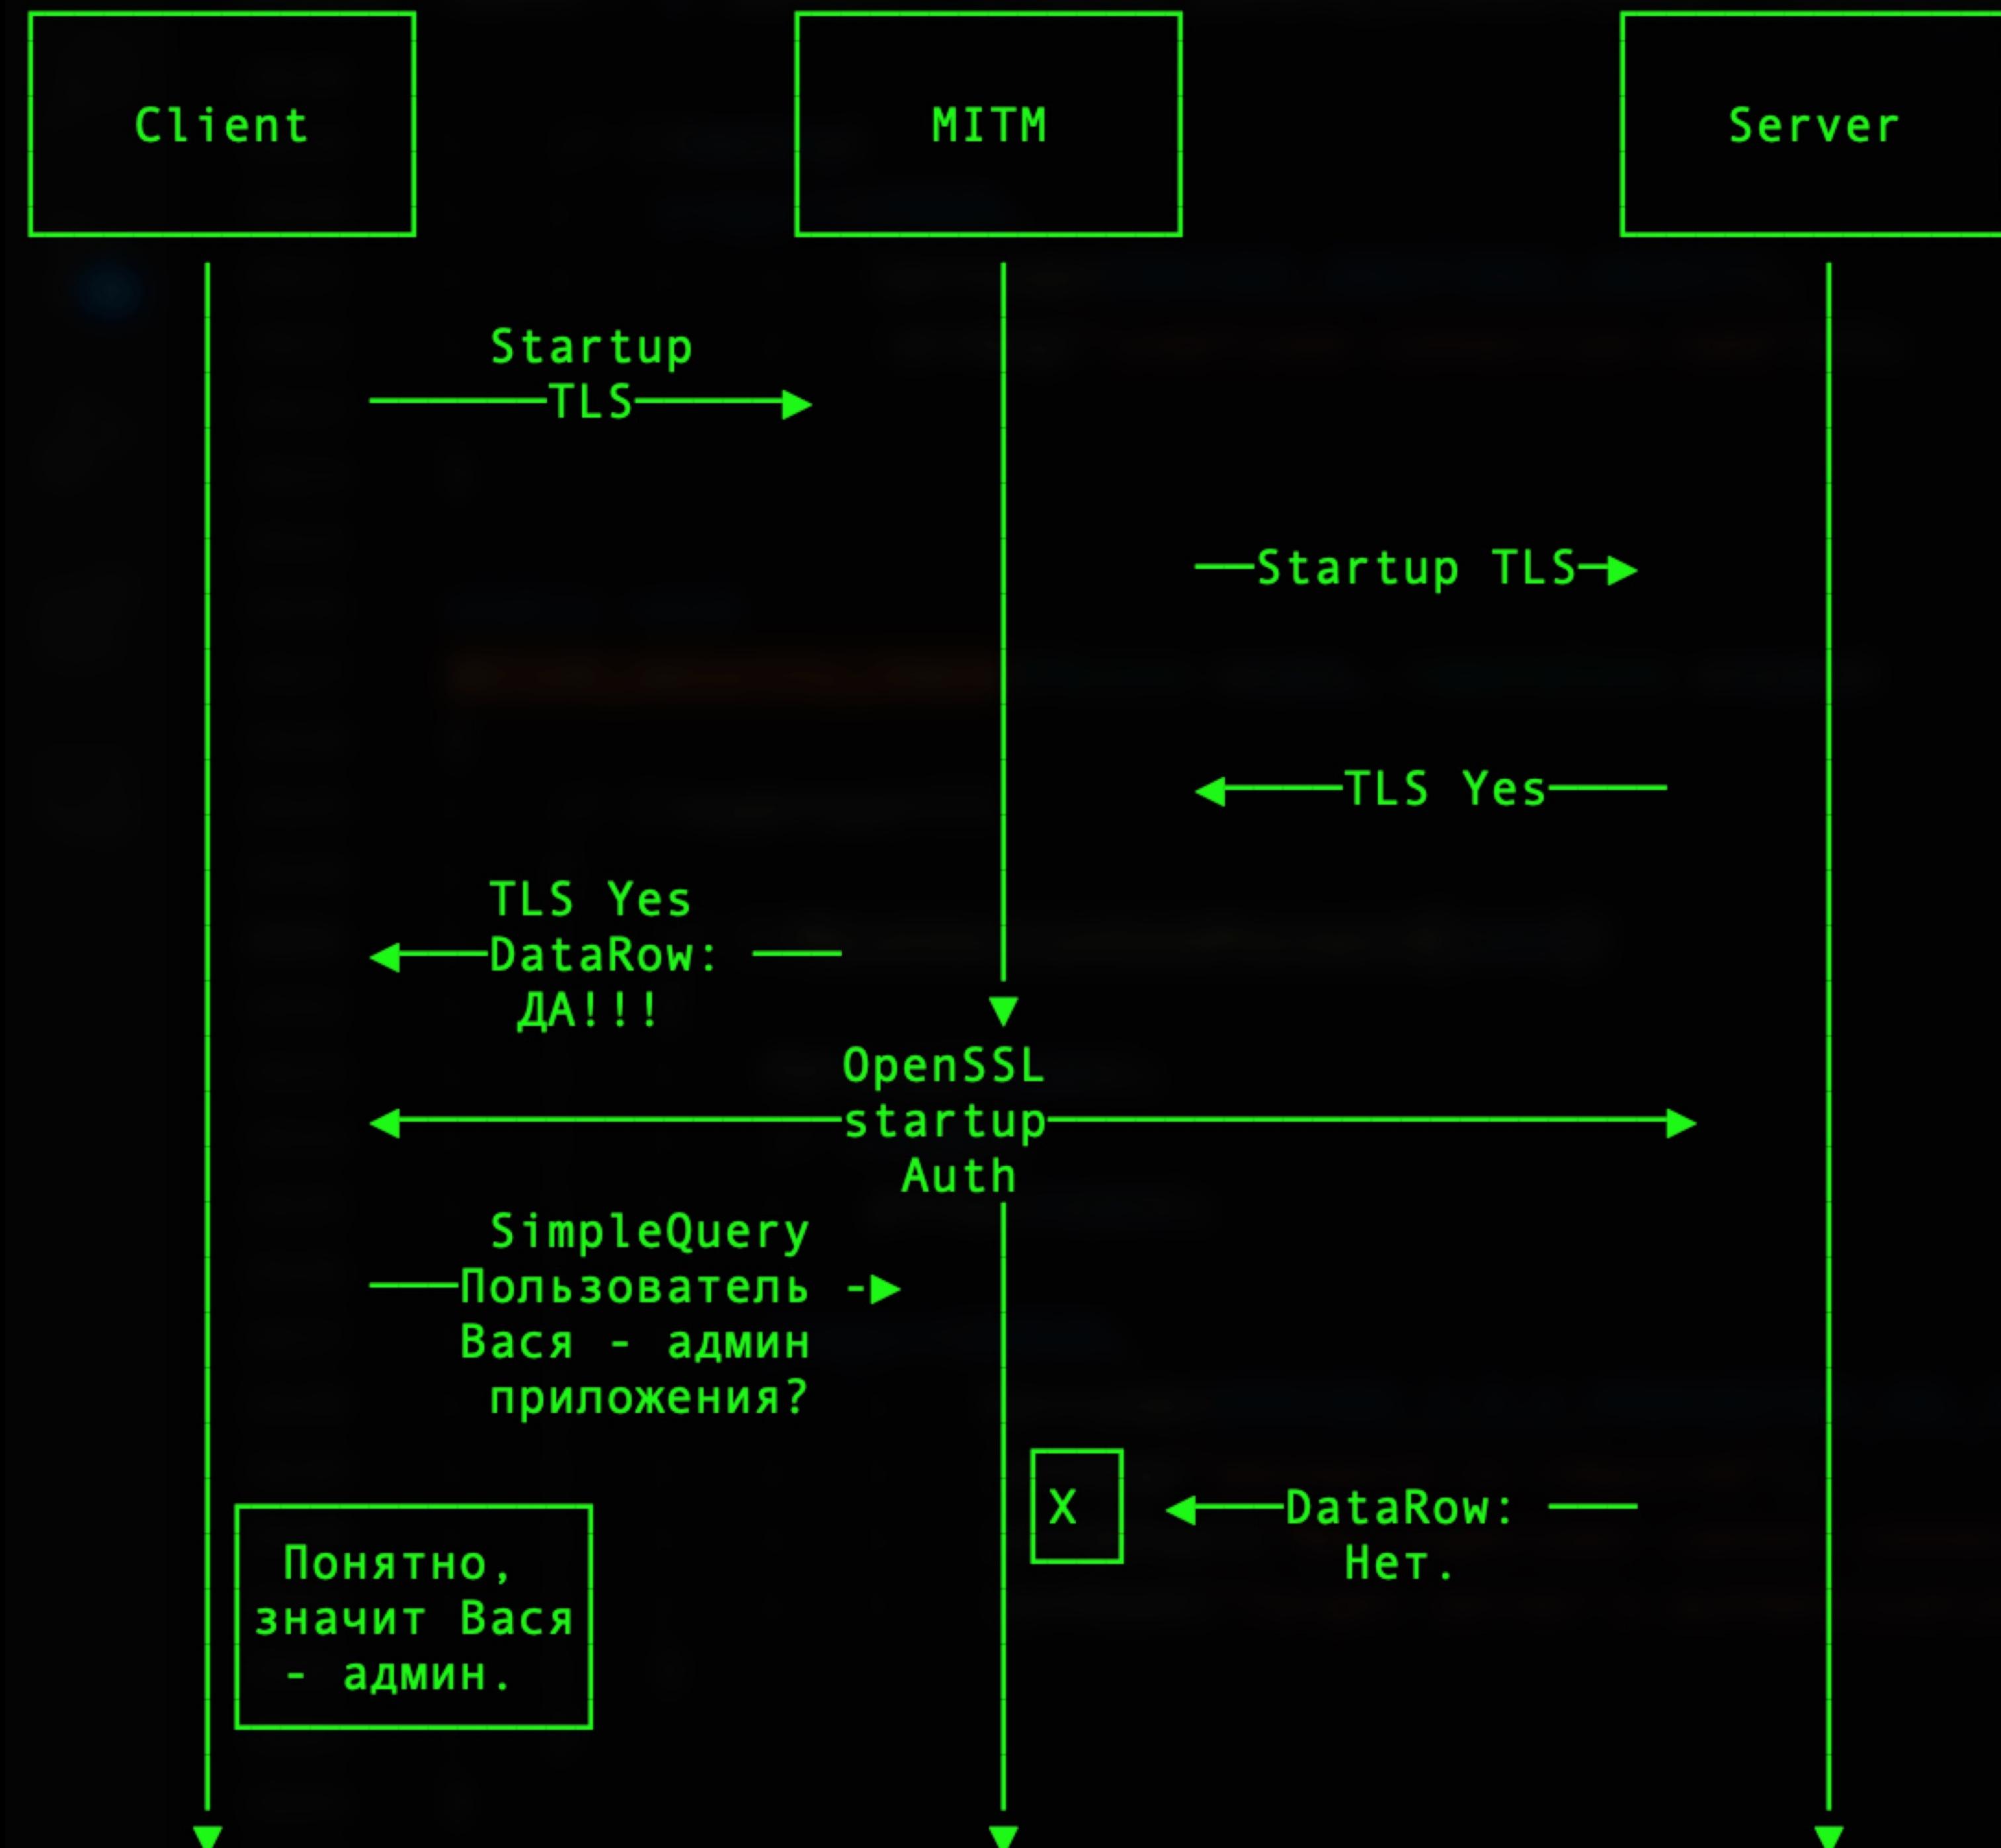
CVE-2021-23214: TLS аутентификация

8.1

Fixed in 14.1, 13.5, 12.9, 11.14, 10.19, 9.6.24 (11 ноября 2021)

X





← → C

https://github.com/postgres/postgres/commit/e52daaak

 Search or jump to... / Pull requests Issues Codespaces Market

postgres / postgres Public

<> Code Pull requests 1 Actions Security Insights

✓ **Reject CancelRequestPacket having unexpected length.**

When the length was too short, the server read outside the allocation. That yielded the same log noise as sending the correct length with (backendPID, cancelAuthCode) matching nothing. Change to a message about the unexpected length. Given the attacker's lack of control over the memory layout and the general lack of diversity in memory layouts at the code in question, we doubt a would-be attacker could cause a segfault. Hence, while the report arrived via security@postgresql.org, this is not a vulnerability. Back-patch to v11 (all supported versions).

Andrey Borodin, reviewed by Tom Lane. Reported by Andrey Borodin.

master

 nmisch committed on Jan 21

▼ ⌂ 7 src/backend/postmaster/postmaster.c ⌂

... @@ -2016,6 +2016,13 @@ ProcessStartupPacket(Port *port, bool ssl_done, bool gss_done)

| | 2016 | 2016 |
|------|------|---|
| 2017 | 2017 | if (proto == CANCEL_REQUEST_CODE) |
| 2018 | 2018 | { |
| 2019 | + | if (len != sizeof (CancelRequestPacket)) |
| 2020 | + | { |
| 2021 | + | ereport (COMMERR0R, |
| 2022 | + | (errcode (ERRCODE_PROTOCOL_VIOLATION), |
| 2023 | + | errmsg ("invalid length of startup packet"))); |
| 2024 | + | return STATUS_ERROR; |
| 2025 | + | } |
| 2019 | 2026 | processCancelRequest (port, buf); |
| 2020 | 2027 | /* Not really an error, but we don't want to proceed further */ |
| 2021 | 2028 | return STATUS_ERROR; |

... ↴



Спасибо!

Андрей Бородин

Postgres Contributor